

# Multiprocessor Real-Time Scheduling

Embedded System Software Design

Prof. Ya-Shu Chen  
National Taiwan University of  
Science and Technology

# Outline

- Multiprocessor Real-Time Scheduling
- Global Scheduling
- Partitioned Scheduling
- Semi-partitioned Scheduling

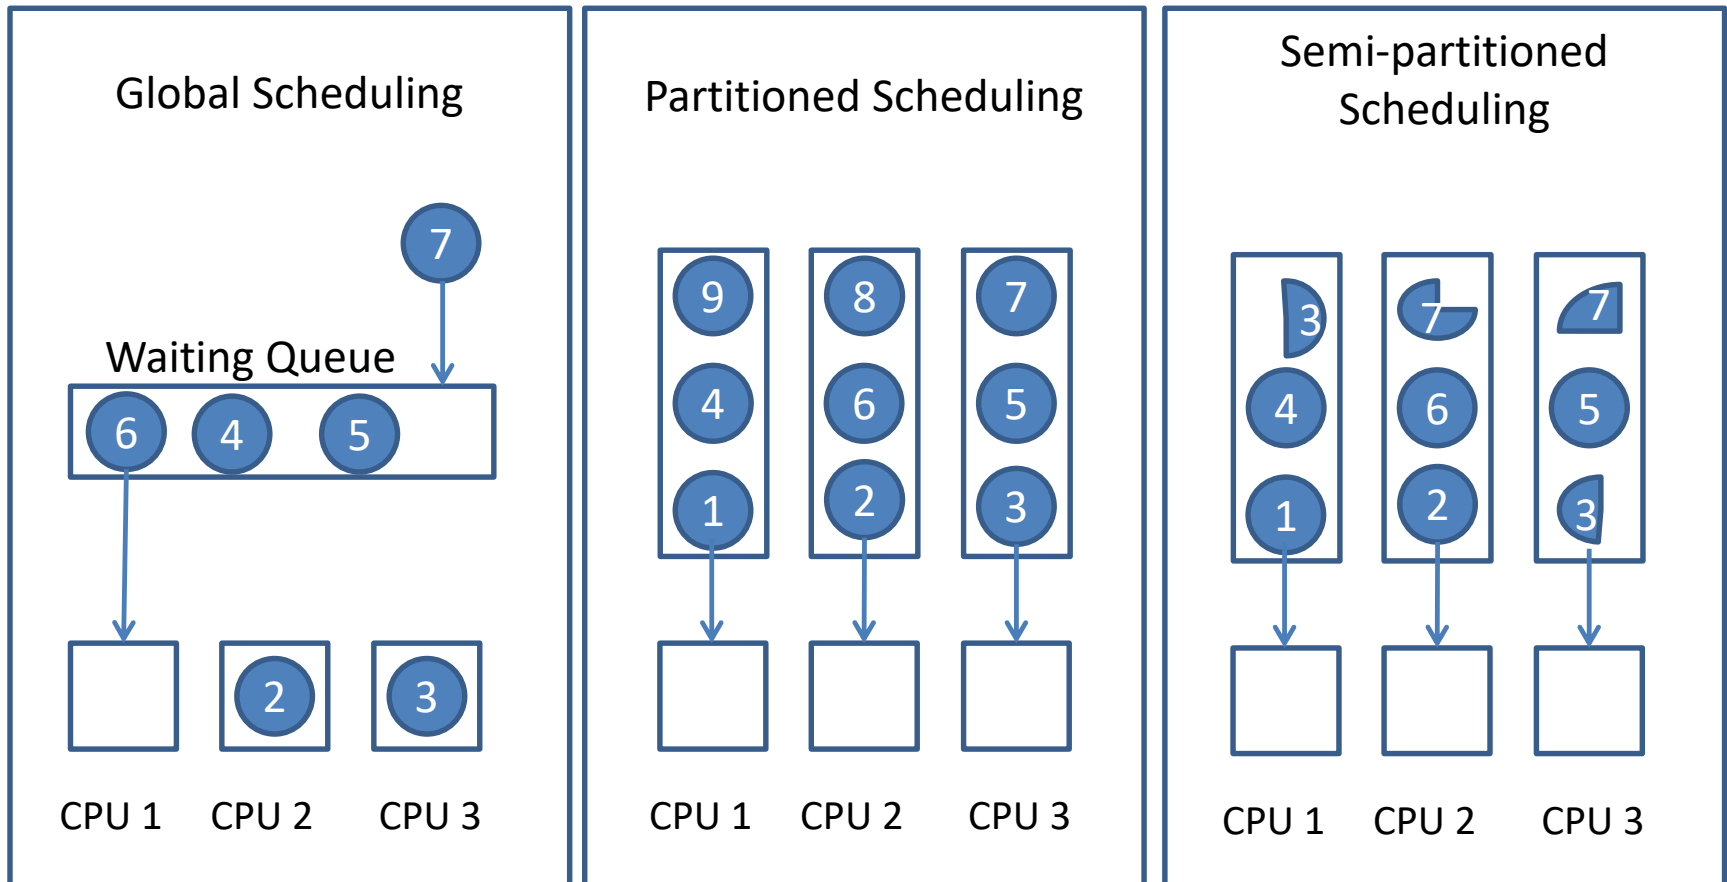
# Multiprocessor Models

- **Identical (Homogeneous):** All the processors have the same characteristics, i.e., the execution time of a job is independent on the processor it is executed.
- **Uniform:** Each processor has its own speed, i.e., the execution time of a job on a processor is proportional to the speed of the processor.
  - A faster processor always executes a job faster than slow processors do.
  - For example, multiprocessors with the same instruction set but with different supply voltages/frequencies.
- **Unrelated (Heterogeneous):** Each job has its own execution time on a specified processor
  - A job might be executed faster on a processor, but other jobs might be slower on that processor.
  - For example, multiprocessors with different instruction sets.

# Scheduling Models

- Global Scheduling:
  - A job may execute on any processor.
  - The system maintains a global ready queue.
  - Execute the  $M$  highest-priority jobs in the ready queue, where  $M$  is the number of processors.
  - It requires high on-line overhead.
- Partitioned Scheduling:
  - Each task is assigned on a dedicated processor.
  - Schedulability is done individually on each processor.
  - It requires no additional on-line overhead.
- Semi-partitioned Scheduling:
  - Adopt task partitioning first and reserve time slots (bandwidths) for tasks that allow migration.
  - It requires some on-line overhead.

# Scheduling Models



# Global Scheduling

- All ready tasks are kept in a global queue
- A job can be migrated to any processor.
- Priority-based global scheduling:
  - Among the jobs in the global queue, the  $M$  highest priority jobs are chosen to be executed on  $M$  processors.
  - Task migration here is assumed with no overhead.
- Global-EDF: When a job finishes or arrives to the global queue, the  $M$  jobs in the queue with the shortest absolute deadlines are chosen to be executed on  $M$  processors.
- Global-RM: When a job finishes or arrives to the global queue, the  $M$  jobs in the queue with the highest priorities are chosen to be executed on  $M$  processors.

# Global Scheduling

- Advantages:
  - Effective utilization of processing resources (if it works)
  - Unused processor time can easily be reclaimed at run-time (mixture of hard and soft RT tasks to optimize resource utilization)
- Disadvantages:
  - Adding processors and reducing computation times and other parameters can actually decrease optimal performance in some scenarios!
  - Poor resource utilization for hard timing constraints
  - Few results from single-processor scheduling can be used

# Schedule Anomaly

- **Anomaly 1**

A decrease in processor demand from higher-priority tasks can *increase* the interference on a lower-priority task because of the change in the time when the tasks execute


- **Anomaly 2**

A decrease in processor demand of a task *negatively* affects the task itself because the change in the task arrival times make it suffer more interference




# Anomaly 1

task set: schedulable



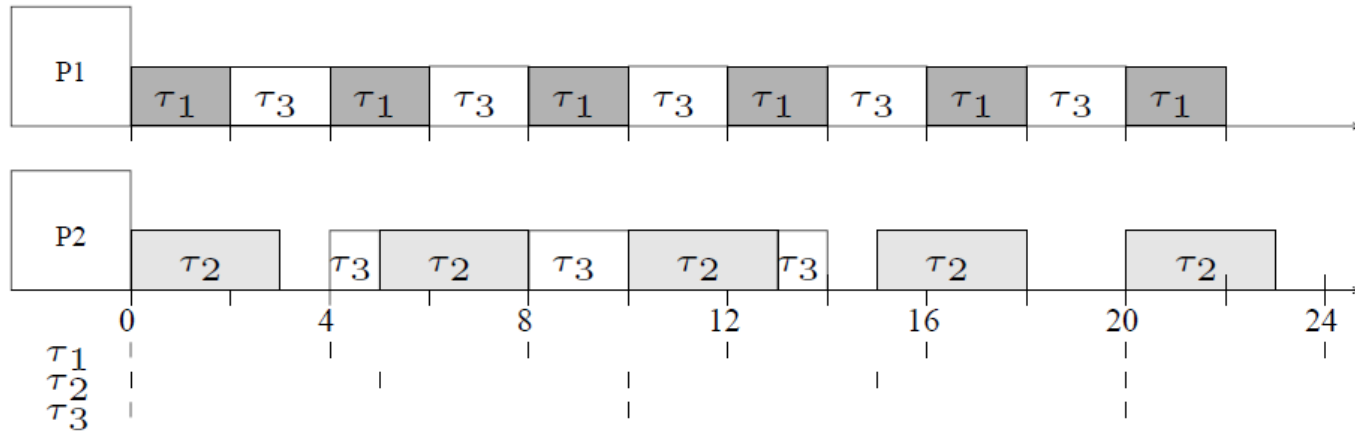
task set: unschedulable



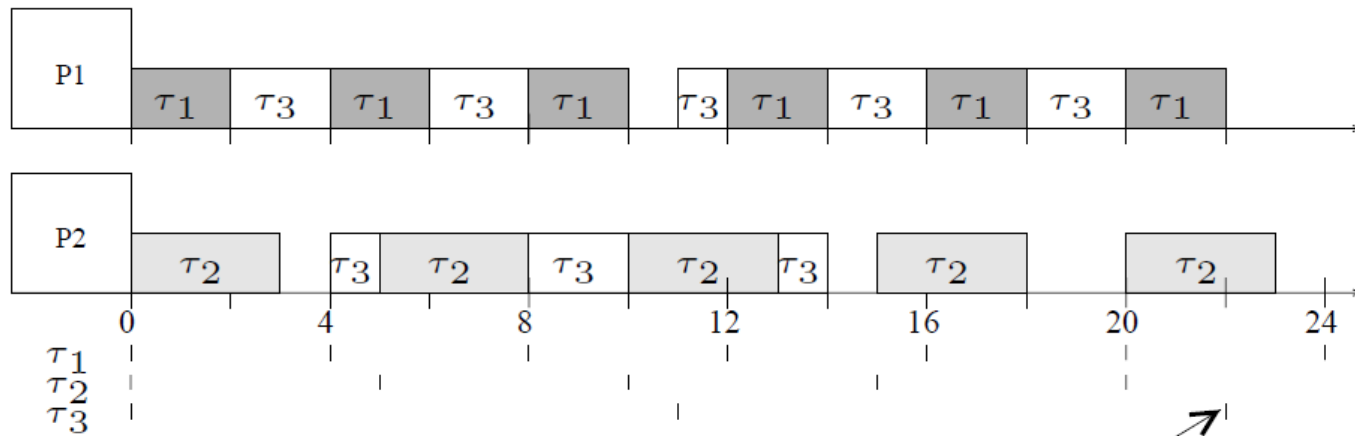
$\tau_3$  needs to execute  
two more time units.

# Anomaly 2

task set: schedulable



task set: unschedulable



$\tau_3$  needs to execute  
one more time unit.

# Dhall effect

- Dhall effect : For Global-EDF or Global-RM, the least upper bound for schedulability analysis is at most 1.
- On 2 processors:

Task	T	D	C	U
T1	10	10	5	0.5
T2	10	10	5	0.5
T3	12	12	8	0.67

- T3 is not schedulable

# Schedulability Test

- A set of periodic tasks  $t_1, t_2, \dots, t_N$  with implicit deadlines is schedulable on  $M$  processors by using preemptive Global EDF scheduling if

$$\sum_{i=1}^N \frac{C_i}{T_i} \leq M(1 - \frac{C_k}{T_k}) + \frac{C_k}{T_k},$$

where  $t_k$  is the task with the largest utilization  $C_k/T_k$

# Weakness of Global Scheduling

- Migration overhead
- Schedule Anomaly

# Partitioned Scheduling

- Two steps:
  - Determine a mapping of tasks to processors
  - Perform run-time single-processor scheduling
- Partitioned with EDF
  - Assign tasks to the processors such that no processor's capacity is exceeded (utilization bounded by 1.0)
  - Schedule each processor using EDF

# Bin-packing Problem

Given a bin size  $V$  and a list  $a_1, \dots, a_n$  of sizes of the items to pack, find an integer  $B$  and a  $B$ -partition  $S_1 \cup \dots \cup S_B$  of  $\{1, \dots, n\}$  such that  $\sum_{i \in S_k} a_i \leq V$ , for all  $k = 1, \dots, B$ .  
A solution is optimal if it has minimal  $B$ .

**The problem is NP-complete !!**

# Bin-packing to Multiprocessor Scheduling

- The problem concerns packing objects of varying sizes in boxes ("bins") with the objective of minimizing number of used boxes.
  - Solutions (Heuristics): First Fit
- Application to multiprocessor systems:
  - Bins are represented by processors and objects by tasks.
  - The decision whether a processor is "full" or not is derived from a utilization-based schedulability test.



# Partitioned Scheduling

- Advantages:
  - Most techniques for single-processor scheduling are also applicable here
- Partitioning of tasks can be automated
  - Solving a bin-packing algorithm
- Disadvantages:
  - Cannot exploit/share all unused processor time
  - May have very low utilization, bounded by 50%

# Partitioned Scheduling Problem

Given a set of tasks with arbitrary deadlines, the objective is to decide a feasible task assignment onto  $M$  processors such that all the tasks meet their timing constraints, where  $C_i$  is the execution time of task  $t_i$  on any processor  $m$ .

# Partitioned Algorithm

- First-Fit: choose the one with the smallest index
- Best-Fit: choose the one with the maximal utilization
- Worst-Fit: choose the one with the minimal utilization

# Partitioned Example

- 0.2 -> 0.6 -> 0.4 -> 0.7 -> 0.1 -> 0.3



First Fit



Best Fit



Worst Fit

# EDF with First Fit

**Input:** A task set  $\tau = \{\tau_1, \tau_2, \dots, \tau_n\}$  and a set of processors  $\{p_1, \dots, p_m\}$

**Output:**  $j$ ; number of processors required.

```
1.   $i := 1; j := 1; k_q = 0; (\forall_q)$ 
2.  while ( $i \leq n$ ) do
3.       $q := 1;$ 
4.      while ( $(U_q + u_i) > 1$ ) do
5.           $q := q + 1; /* \text{increase the processor index} */$ 
6.           $U_q := U_q + u_i; k_q := k_q + 1;$ 
7.          if ( $q > j$ ) then
8.               $j := q;$ 
9.           $i := i + 1;$ 
10. return ( $j$ );
11. end
```

# Schedulability Test

Lopez [3] proves that the **worst-case achievable utilization** for EDF scheduling and FF allocation (EDF-FF) takes the value

If all the tasks have an utilization factor C/T under a value  $\alpha$ , where  $m$  is the number of processors

$$U_{wc}^{EDF-FF}(m, \beta) = \frac{\beta m + 1}{\beta + 1} \quad \text{where} \quad \beta = \lfloor 1/\alpha \rfloor$$


# Demand Bound Function

- Define demand bound function  $dbf(\tau_i, t)$  as

$$dbf(\tau_i, t) = \max \left\{ 0, \left\lfloor \frac{t + T_i - D_i}{T_i} \right\rfloor \right\} C_i = \max \left\{ 0, \left\lfloor \frac{t - D_i}{T_i} \right\rfloor + 1 \right\} C_i.$$

- We need approximation to enforce polynomial-time schedulability test

$$dbf^*(\tau_i, t) = \begin{cases} 0 & \text{if } t < D_i \\ (\frac{t-D_i}{T_i} + 1) C_i & \text{otherwise.} \end{cases}$$



# Deadline Monotonic Partition

**Input:**  $\mathbf{T}, M$ ;

- 1: re-index (sort) tasks such that  $D_i \leq D_j$  for  $i < j$ ;
- 2:  $\mathbf{T}_i \leftarrow \emptyset, U_i \leftarrow 0, \forall m = 1, 2, \dots, M$ ;
- 3: **for**  $i = 1$  to  $N$ , where  $N = |\mathbf{T}|$  **do**
- 4:     **for**  $m = 1$  to  $M$  **do**
- 5:         **if**  $\frac{C_i}{T_i} + \sum_{\tau_j \in \mathbf{T}_m} \frac{C_j}{T_j} \leq 1$  and  $C_i + \sum_{\tau_j \in \mathbf{T}_m} dbf^*(\tau_j, D_i) \leq D_i$  **then**
- 6:             assign task  $\tau_i$  onto processor  $m$  and  $\mathbf{T}_m \leftarrow \mathbf{T}_m \cup \{\tau_i\}$ ;
- 7:             break;
- 8:     **if**  $\tau_i$  is not assigned **then**
- 9:         return "The task assignment fails";
- 10: return feasible task assignment  $\mathbf{T}_1, \mathbf{T}_2, \dots, \mathbf{T}_M$ ;



# Schedulability Test

**Theorem 4** *Any sporadic task system  $\tau$  is successfully scheduled by Algorithm PARTITION on  $m$  unit-capacity processors, for any*

$$m \geq \left( \frac{2\delta_{\text{sum}} - \delta_{\text{max}}}{1 - \delta_{\text{max}}} + \frac{u_{\text{sum}} - u_{\text{max}}}{1 - u_{\text{max}}} \right) \quad (14)$$

$$\begin{array}{ll} \delta_{\text{max}} & \stackrel{\text{def}}{=} \max_{i=1}^n (e_i/d_i) & u_{\text{max}} & \stackrel{\text{def}}{=} \max_{i=1}^n (u_i) \\ \delta_{\text{sum}} & \stackrel{\text{def}}{=} \max_{t>0} \left( \frac{\sum_{j=1}^n \text{DBF}(\tau_j, t)}{t} \right) & u_{\text{sum}} & \stackrel{\text{def}}{=} \sum_{j=1}^n u_j \end{array}$$

# Weakness of Partitioned Scheduling



- Restricting a task on a processor reduces the schedulability
- Restricting a task on a processor makes the problem NP-hard
- Example: Suppose that there are  $M$  processors and  $M + 1$  tasks with the same period  $T$  and the (worst-case) execution times of all these  $M + 1$  tasks are  $T/2 + e$  with  $e > 0$ 
  - With partitioned scheduling, it is not schedulable

# Semi-partitioned Scheduling

- Tasks are first partitioned into processor.
- To reduce the utilization, we again pick the processor with the minimum task utilization
- If a task cannot fit into the picked processor, we will have to split it into multiple (two or even more) parts.
- If  $t_i$  is split and assigned to a processor  $m$  and the utilization on processor  $m$  after assigning  $t_i$  is at most  $U(\text{scheduler}, N)$ , then  $t_i$  is so far schedulable.


# Semi-partitioned EDF

- $T_{\min}$  is the minimum period among all the tasks.
- By a user-designed parameter  $k$ , we divide time into slots with length  $S = T_{\min}/k$ .
- We can use the first-fit approach by splitting a task into 2 subtasks, in which one is executed on processor  $m$  and the other is executed on processor  $m + 1$ .
- Execution of a split task is only possible in the reserved time window in the time slot.
- Applying first-fit algorithm, by taking SEP as the upper bound of utilization on a processor.
- If a task does not fit, split this task into two subtasks and allocate a new processor, one is assigned on the processor under consideration, and the other is assigned on the newly allocated processor.




# Semi-partitioned EDF

- For each time slot, we will reserve two parts.



If a task  $t_i$  is split, the task can be served only within these two pre-defined time slots with length  $x_i$  and  $y_i$ .



A processor can host two split tasks,  $t_i$  and  $t_j$ .  $t_i$  is served at the beginning of the time slot, and  $t_j$  is served at the end.

The schedule is EDF, but if a split task instance is in the ready queue, it is executed in the reserved time region.

# Semi-partitioned EDF

- We can assign all the tasks  $t_i$  with  $U_i > SEP$  on a dedicated processor. So, we only consider tasks with  $U_i$  no larger  $SEP$ .

```
1:  $m \leftarrow 1, U_m \leftarrow 0;$ 
2: for  $i = 1$  to  $N$ , where  $N = |\mathbf{T}|$  do
3:   if  $\frac{C_i}{T_i} + U_m \leq SEP$  then
4:     assign task  $\tau_i$  on processor  $m$ ;
5:      $U_m \leftarrow U_m + \frac{C_i}{T_i};$ 
6:   else
7:     assign task  $\tau_i$  on processor  $m$  with  $lo\_split(\tau_i)$  set to  $SEP - U_m$  and on
       processor  $m + 1$  with  $high\_split(\tau_i)$  set to  $\frac{C_i}{T_i} - (SEP - U_m);$ 
8:      $m \leftarrow m + 1$  and  $U_m \leftarrow \frac{C_i}{T_i} - (SEP - U_m);$ 
```

When executing, the reservation to serve  $t_i$  is to set  $x_i$  to  $S \times (f + lo\_split(t_i))$  and  $y_i$  to  $S \times (f + high\_split(t_i))$ .  $SEP$  is set as a constant.

# Two Split Tasks on a Processor

- For split tasks to be schedulable, the following sufficient conditions have to be satisfied
  - $\text{lo\_split}(t_i) + f + \text{high\_split}(t_i) + f \leq 1$  for any split task  $t_i$ .
  - $\text{lo\_split}(t_j) + f + \text{high\_split}(t_i) + f \leq 1$  when  $t_i$  and  $t_j$  are assigned on the same processor.
- Therefore, the “magic value” SEP

$$SEP \leq 1 - 2f \leq 1 - 2(\sqrt[2]{\kappa(\kappa + 1)} - \kappa).$$

- However, we still have to guarantee the schedulability of the non-split tasks. It can be shown that the sufficient condition is

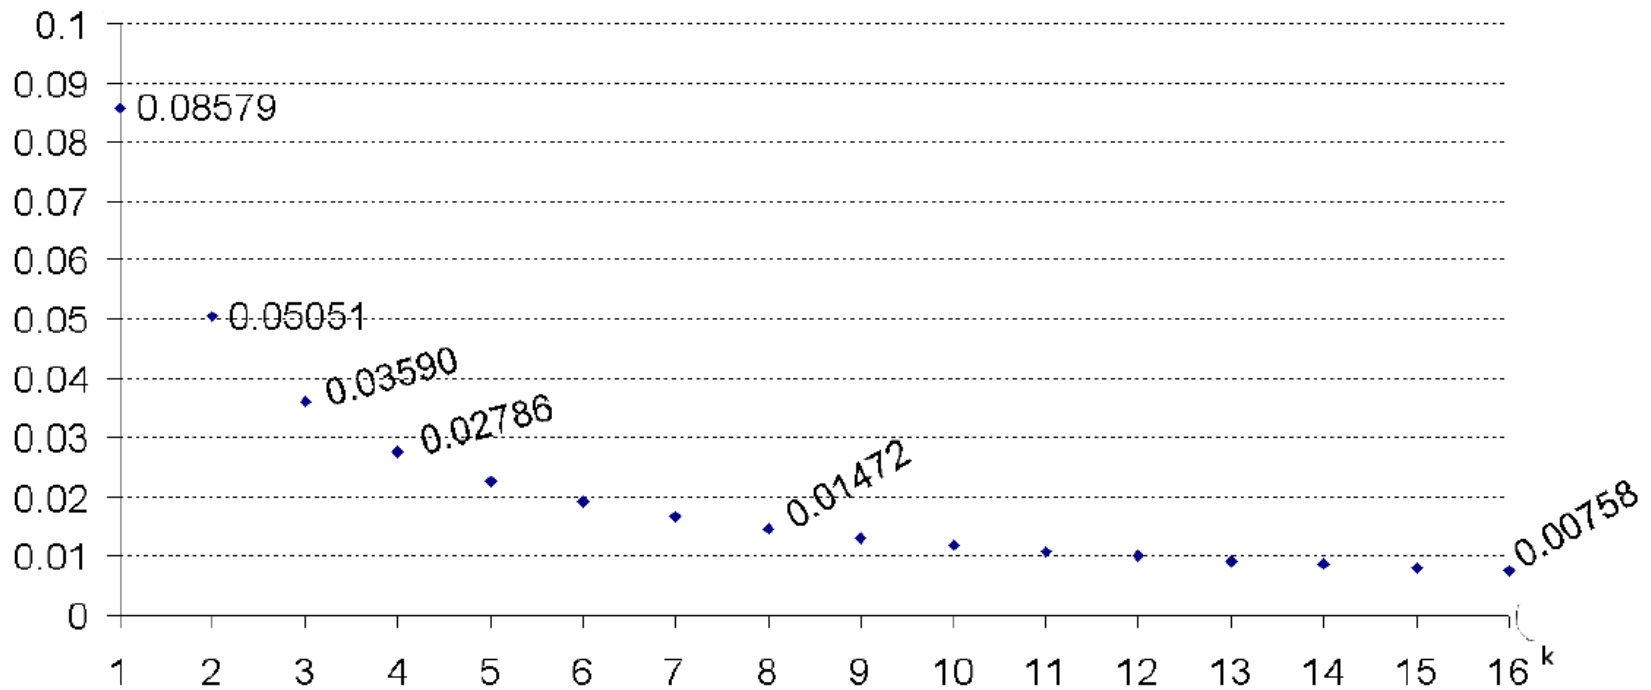
$$SEP \leq 1 - 4f \leq 1 - 4(\sqrt[2]{\kappa(\kappa + 1)} - \kappa).$$



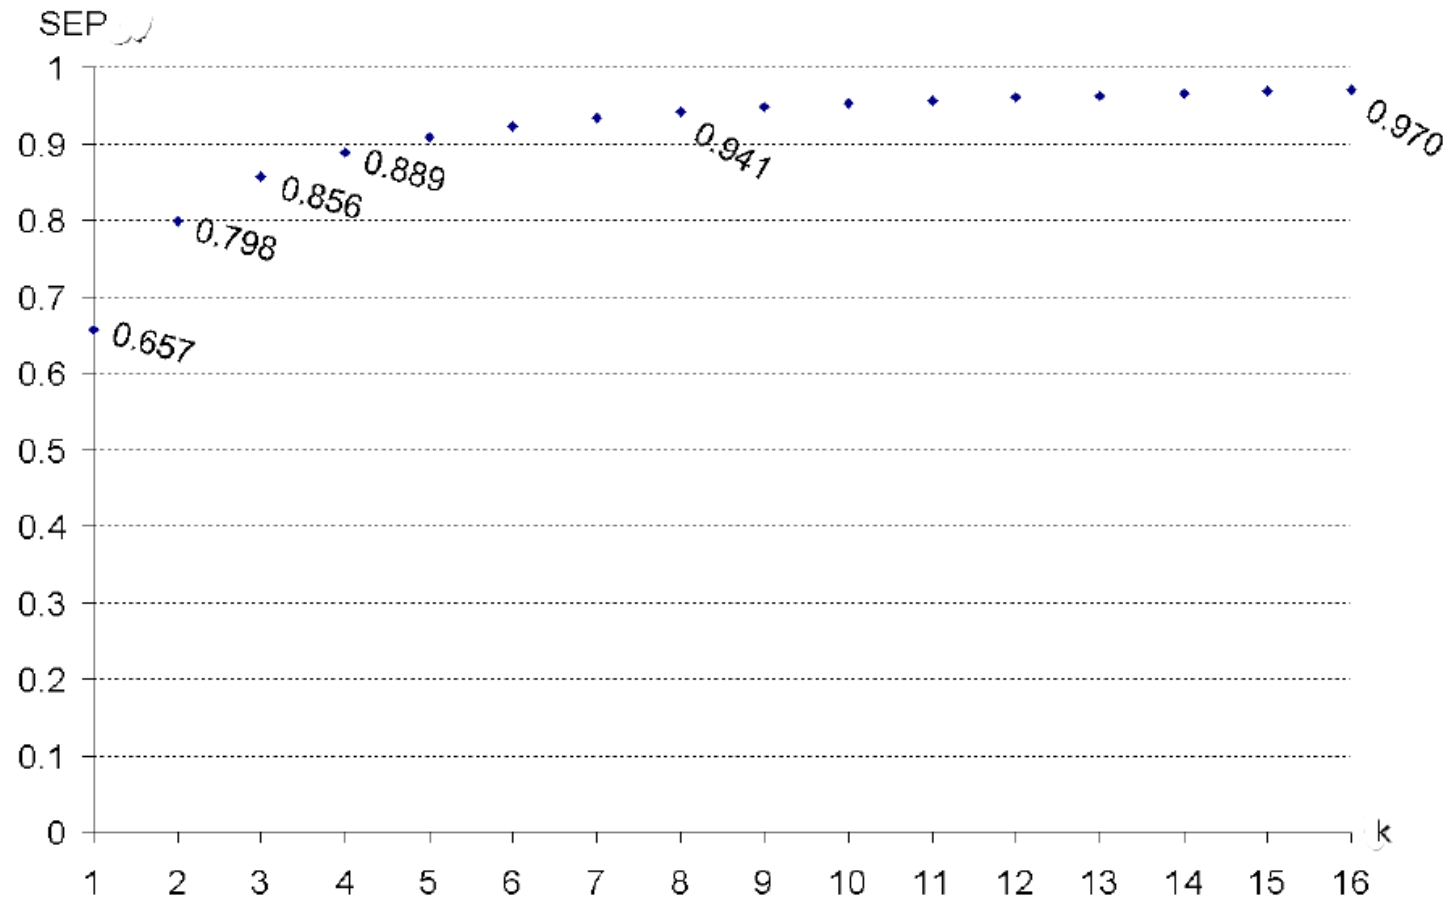
# Schedulability Test

By taking  $SEP$  as  $1 - 4(\sqrt[2]{\kappa(\kappa + 1)} - \kappa)$  and  $f = \sqrt[2]{\kappa(\kappa + 1)} - \kappa$ , the above algorithm guarantees to derive feasible schedule if  $\sum_{\tau_i \in \mathbf{T}} \frac{C_i}{T_i} \leq M' \cdot SEP$  and  $\frac{C_i}{T_i} \leq SEP$  for all tasks  $\tau_i$ .

# Magic Values: f



# Magic Values: SEP



# Reference

- Multiprocessor Real-Time Scheduling
  - Dr. Jian-Jia Chen: Multiprocessor Scheduling. Karlsruhe Institute of Technology (KIT): 2011-2012
- Global Scheduling
  - Sanjoy K. Baruah: Techniques for Multiprocessor Global Schedulability Analysis. RTSS 2007: 119-128
- Partitioned Scheduling
  - Sanjoy K. Baruah, Nathan Fisher: The Partitioned Multiprocessor Scheduling of Sporadic Task Systems. RTSS 2005: 321-329
- Semi-partitioned Scheduling
  - Björn Andersson, Konstantinos Bletsas: Sporadic Multiprocessor Scheduling with Few Preemptions. ECRTS 2008: 243-252

See You Next Week