

Real-Time End-to-End Scheduling

Embedded System Software Design

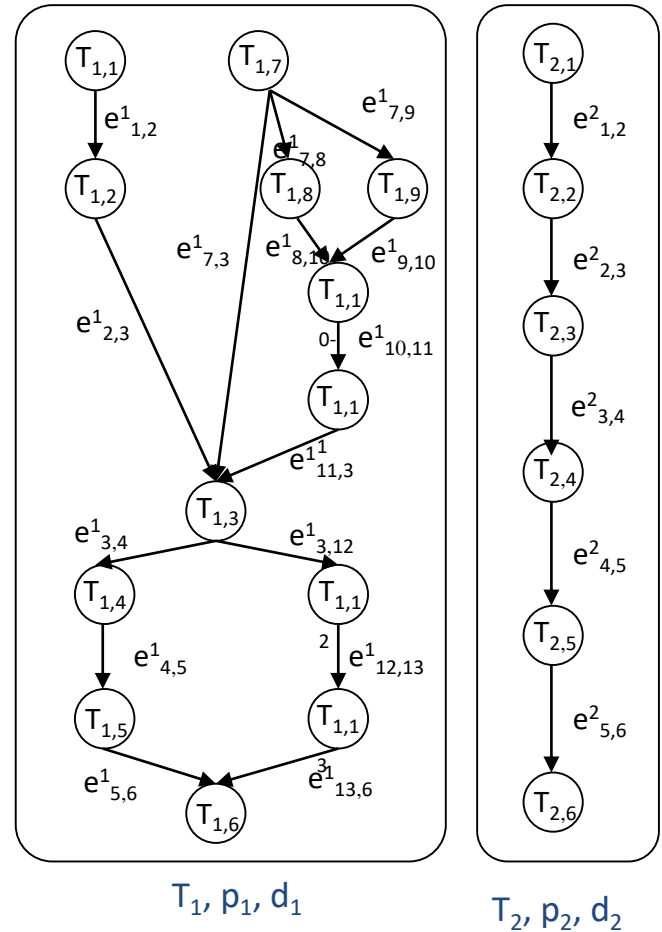

Prof. Ya-Shu Chen
National Taiwan University of
Science and Technology

End-to-End Scheduling

- Task model
 - Each task needs to execute on a set of processors in a certain order
 - Each task may require a different order
- Problems in End-to-End scheduling
 - Priority assignment
 - Assign fixed priorities to tasks so that the system is schedulable
 - Synchronization of tasks
 - Control the releases of subtask instances (non-first subtasks)
 - Schedulability analysis
 - For a given priority assignment and a given synchronization protocol, whether every instance of each task meets its deadline

System Model

- Platform: A set of processors
- Task graph, $G = \{T_1, T_2, \dots, T_n\}$
 - Sink node
 - Deadline: d_i
 - Precedence edge: $e_{j,k}^i$
 - Predecessors and Successors
 - Period or Minimum separation time: p_i
- Characteristics of $T_{i,j}$:
 - Execution time (on processor m): $c_{i,j}^m$



Priority Assignment

- To find feasible priority assignments off-line if all tasks executed are known in prior
- NP-hard problem
- Algorithms
 - Branch and bound
 - Search algorithm
 - Simulated annealing
 - Generic algorithm
 - Heuristic
 - Deadline assignment

Deadline Assignment

- Ultimate deadline
 - $UD_{i,k} = D_i$
- Effective deadline
 - $ED_{i,k} = D_i - \sum_{l=k+1}^{n(i)} e_{i,l}$
- Proportional deadline
 - $PD_{i,k} = D_i e_{i,k} / e_i$
- Normalized Proportional deadline
 - $NPD_{i,k} = D_i \frac{e_{i,k} U(V_{i,k})}{\sum_{l=1}^{n(i)} e_{i,l} U(V_{i,l})}$
 - $U(V_{i,l})$ is the total utilization of the all the subtasks that execute on the processor $V_{i,l}$

Example

$T_{i,k}$	$V_{i,k}$	p_i	$e_{i,k}$	$UD_{i,k}$	$ED_{i,k}$	$PD_{i,k}$	$NPD_{i,k}$
$T_{1,1}$	P_1	15	1	15	11	3	2.0
$T_{1,3}$	P_1	15	2	15	15	6	4.1
$T_{2,1}$	P_1	20	4	20	20	20	20
$T_{3,1}$	P_2	2	1	2	2	2	2
$T_{1,2}$	P_2	15	2	15	13	6	8.9
$T_{4,1}$	P_2	20	5	20	20	20	20


The Synchronization Problem

- Given that
 - Priorities are assigned to subtasks in a task chain using some fixed priority assignment algorithm
- How do we coordinate the release of subtasks in a task chain so that
 - Precedence constraints among subtasks are satisfied
 - Subtask deadlines are met
 - End-to-end deadlines are met

Synchronization Protocols

- Direct Synchronization (DS) Protocol
 - Simple and straightforward
- Phase Modification (PM) Protocol
 - Used by flow-shop tasks
 - Extension called Modified Phase Modification (MPM) Protocol
- Release Guard Protocol
 - Reclaim the idle time

Example



$T_{i,j}$ – j^{th} subtask of task T_i

(period, execution time)


Period = relative deadline of parent task

Task T_3 releases at 6


Direct Synchronization Protocol

- Greedy strategy
- On completion of subtask
 - A synchronization signal sent to the next processor
 - Successor subtask competes with other tasks/subtasks on the next processor

Greedy Example



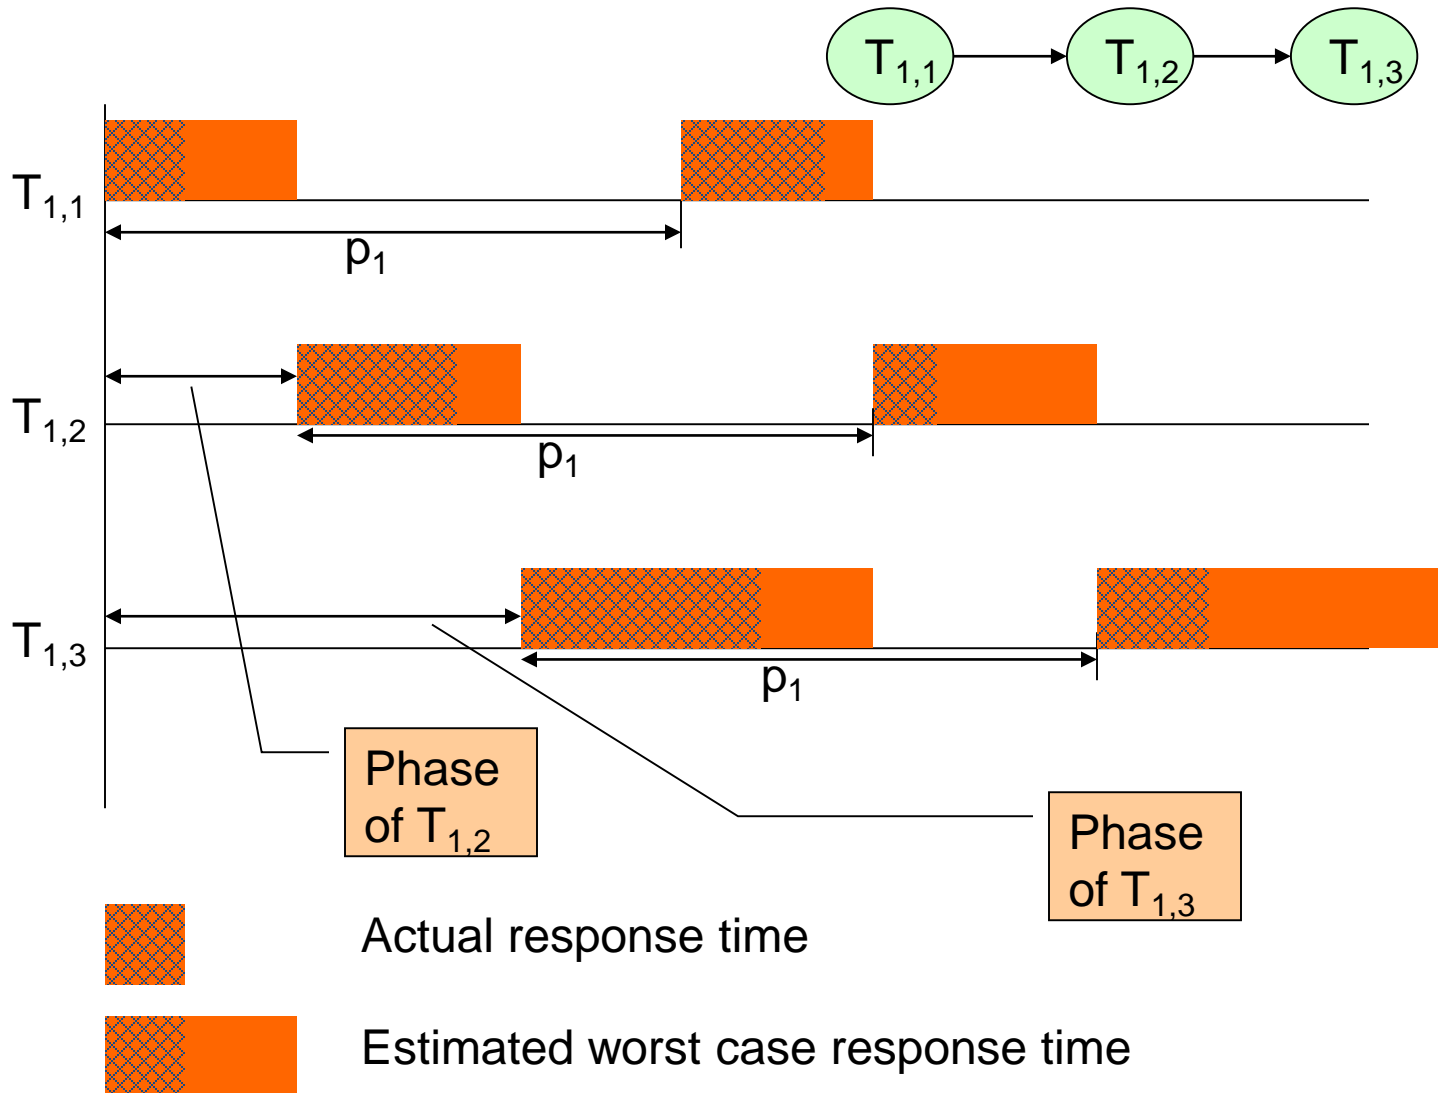
Non-greedy Example




Phase Modification Protocol

- Release subtasks periodically
 - According to the periods of their parent tasks
- Each subtask given its own phase
- Phase determined by subtask precedence constraints

Phase Modification Protocol (1/2)




Phase Modification Protocol (2/2)




Phase Modification Protocol - Analysis

- Periodic timer interrupt to release subtasks
- Centralized clock or strict clock synchronization
- Task overruns could cause precedence constraint violations

Modified PM Protocol (1/2)




Modified PM Protocol (2/2)




Release Guard Protocol

- A guard variable – *release guard* - associated with each subtask
- Release guard used to control release of each subtask
 - Contains next release time of subtask
- Synchronization signals as MPM
- Release guard updated
 - On getting synchronization signal
 - During idle time

Release Guard Protocol



Release Guard Protocol



Release Guard Protocol - Analysis

- Shares the same advantages as MPM
- Upper bound on EER still the same as MPM
 - Since upper bound on release time enforced by release guard
- Lower bound on EER less than that of MPM
 - If there are idle times
 - Results in lower average EER (end-to-end response time)

Schedulability Analysis

An upper bound W_i to the end-to-end response time of any periodic task T_i in a fixed-priority system synchronized according to the MPM protocol or the RG protocol is given by

$$W_i = \sum_{k=1}^{n(i)} W_{i,k}$$

and

$$W_{i,k} = \frac{e_{i,k} + b_{i,k} + \sum_{\phi_{j,l} \leq \phi_{i,k} \text{ and } \tau_{j,l} \in V_{i,k}} e_{j,l}}{1 - \sum_{\phi_{j,l} < \phi_{i,k} \text{ and } \tau_{j,l} \in V_{i,k}} u_{j,l}}$$


where $n(i)$ is the number of subtasks in T_i , $\phi_{i,k}$ is the priority of $\tau_{i,k}$, and the upper bound $W_{i,k}$ to the response time of every subtask $T_{i,k}$ is obtained by considering only subtasks on the same processor $V_{i,k}$, and by treating every such subtask $T_{j,l}$ as periodic task whose period is equal to the period p_j of the parent task T_j .

General Scheduling Test (GST)

- Response time analysis
 - The response time of the job of T_i at critical instant can be calculated by the following recursive function

$$r_0 = \sum_{\forall i} c_i$$

$$r_n = \sum_{\forall i} c_i \left\lceil \frac{r_{n-1}}{p_i} \right\rceil$$



- Observation: the sequence of r_x , $x \geq 0$ may or may not converge


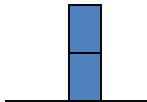
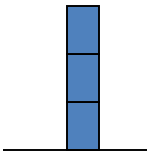

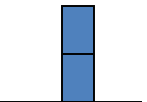
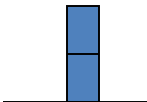

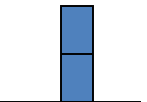

General Scheduling Test (GST)

- Example: $T1=(2,5)$, $T2=(2,7)$, $T3=(3,8)$
 - T1:
 - $R_0=2 \leq 5$ ok
 - T2:
 - $R_0=2+2=4 \leq 7$
 - $R_1=2 * \lceil 4/5 \rceil + 2 * \lceil 4/7 \rceil = 4 \leq 7$ ok
 - T3:
 - $R_0=2+2+3=7 \leq 8$
 - $R_1= 2 * \lceil 7/5 \rceil + 2 * \lceil 7/7 \rceil + 3 * \lceil 7/8 \rceil = 9 > 8$ failed
 - Note: each task succeeds \rightarrow the task set succeeds

Example

$T_{i,k}$	$V_{i,k}$	p_i	$e_{i,k}$	$UD_{i,k}$	$b_{i,k}$	$W_{i,k}$	$W_{i,k}(GST)$
$T_{1,1}$	P_1	15	1	15	0	3	3
$T_{1,3}$	P_1	15	2	15	1	4	3(4)
$T_{2,1}$	P_1	20	4	20	0	8.75	7
$T_{3,1}$	P_2	2	1	2	0	1	1
$T_{1,2}$	P_2	15	2	15	1	8	4(6)
$T_{4,1}$	P_2	20	5	20	0	21.8	14

Comparison of Protocols

	DS	PM	MPM	RG
Implementation complexity	Synch interrupts	Timer interrupts clock synchronization	Synch & timer interrupts	Synch & timer interrupts
Run-time overhead				
Average EER				
Estimated worst case EER				
Inherently missed deadlines	Yes	No		

Reference

- Real-time Systems, Jane Liu
- Bettati, R., ``End-to-end scheduling to meet deadlines in distributed systems,” Ph.D. thesis, University of Illinois at Urbana-Champaign
- Sun, J., ``Fixed-Priority Scheduling of Periodic Tasks With End-to-End Deadlines,” Ph.D. thesis, University of Illinois at Urbana-Champaign